

Perceptual Losses for Real-Time Style Transfer

Aakef Waris, Ali Issaoui



Table of Contents

1. Problem statement
2. The Dataset
3. Network Architecture
4. The Image Transformation Network
5. Training and Loss
6. Results
7. Problems and Further Improvements

Problem:

- Implement style transfer using perceptual loss functions
- Perceptual Loss Functions?
 - Instead of traditional pixel-by-pixel loss, perceptual loss functions are computed from intermediary layers from pre-trained networks
 - This captures higher level feature similarities across images instead of pixel values
 - Effective when trying to capture and recognize features like edges, texture, sharpness, color, etc.
- Compute style AND feature loss functions to converge using Adam optimizer

The Dataset

MSCOCO :




Styles:

COCO is a large-scale object detection, segmentation, and captioning dataset.

Original dataset:

- 330K images**
- 1.5 million object instances**
- 80 object categories**
- 91 stuff categories**

Our subset: **30K images**



Network Architecture

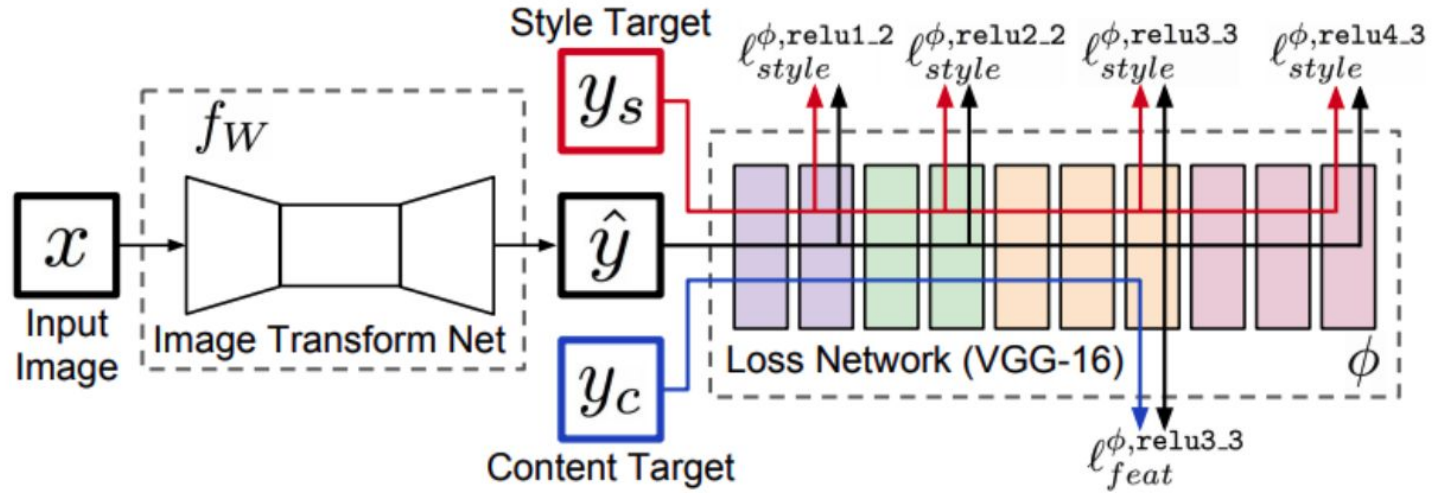


Image Transformation Network

- Initial Down Sampling Layers from 3 to 32 to 64 to 128 channels
- Body of 5 Residual Blocks at 128 channels
- Final Up Sampling Layers from 128 to 64 to 32 to 3 channels for the output image

```
class ImageTransformationNN(torch.nn.Module):  
  
    def __init__(self):  
        super(ImageTransformationNN, self).__init__()  
        self.down_sample = DownSampleConv()  
        self.res = ResidualNet()  
        self.up_sample = UpSampleConv()  
  
    def forward(self, X):  
  
        X = self.down_sample(X)  
        X = self.res(X)  
        y = self.up_sample(X)  
        return y
```

Down/Up -Sampling

- 3 Convolutional Layers
- Used Reflection Padding
- Followed by Instance Normalization

Instance Normalization




Figure 3:




Figure 4:

Reflection Padding

3	5	1
3	6	1
4	7	9

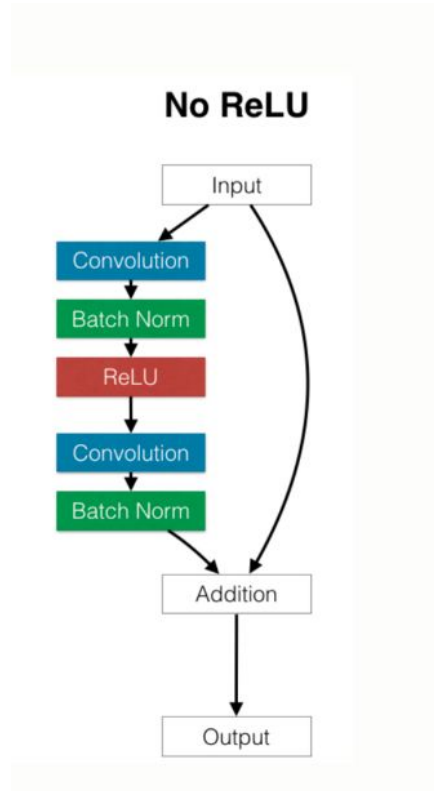
No padding

1	6	3	6	1	6	3
1	5	3	5	1	5	3
1	6	3	6	1	6	3
9	7	4	7	9	7	4
1	6	3	6	1	6	3


(1, 2) reflection padding

Figure 5:

Residual Blocks



VGG-16 Network (Loss Network)



Training and Loss

Feature reconstruction loss (relu3_3):


$$l_{feat}^{\theta,j}(\hat{y}, y) = \frac{1}{C_j H_j W_j} \|\theta_j(\hat{y}) - \theta_j(y)\|_2^2$$

Style Reconstruction Loss(relu1_2, relu2_2, relu3_3, relu4_3):

$$l_{style}^{\theta,j}(\hat{y}, y) = \|G_j^{\theta}(\hat{y}) - G_j^{\theta}(y)\|_F^2$$

Training parameters:

- Epochs: 2
- Batch size: 4
- Optimizer: Adam
- Learning rate: 1e-3



Results

Style Image:



Content Image:



Output Image:



Problems and Further improvements

Problems:

1. Limited computational power
2. Limited time
3. Limited space for the entire COCO dataset (18 Gb)

Further improvements:

1. Continue the training for the Style Transfer task
2. Generalize the perceptual loss method on other transformation tasks:
 - a. Image super-resolution
 - b. Image colorization, and others

References

- The original paper:
<https://arxiv.org/pdf/1603.08155.pdf>
- Instance Norm:
<https://becominghuman.ai/all-about-normalization-6ea79e70894b#:~:text=In%20%E2%80%9CInstance%20Normalization%E2%80%9D,%20mean,sample%20across%20both%20spatial%20dimensions.>
- Reflection Padding:
<https://www.machinecurve.com/index.php/2020/02/10/using-constant-padding-reflection-padding-and-replication-padding-with-keras/>
- Residual Block Architecture:
<http://torch.ch/blog/2016/02/04/resnets.html>

Thank you